

Stochastic Optimization Assignment Report

Giacomo Maino s338682, Mattia Molinari s337194, Alessandro Perlo s337131

February 21st, 2025

Abstract

This report provides a comprehensive overview of methodologies and decisions taken during the development of a metaheuristic algorithm to solve the Integrated Healthcare Timetabling Competition 2024. The proposal is based on the Tabu Search algorithm, which has been opportunely adapted to the problem at hand.

1 Problem Overview

Integrated healthcare scheduling involves the coordination of resources across multiple services within a unified healthcare system. Its primary goal is to streamline and optimize patient flow throughout various departments and facilities within the hospital. The advantages of integrated healthcare optimization are numerous, including an enhanced patient experience, increased operational efficiency, and better utilization of resources across the entire hospital network.

The *Integrated Healthcare Timetabling Problem* (IHTP) [1] transpose the aforementioned concept into a combinatorial optimization problem, bringing together three distinct operational challenges: Patient-to-Room Assignment (PRA), Nurse-to-Patient Assignment (NPA) and Surgical Case Planning (SCP). More precisely, the IHTP requires to establish (i) the date of admission for patients, (ii) the room where they are admitted and (iii) operating theater where they are operated, (iii) the room nurses are assigned to during each shift. The problem is subject to several constraints, specific to a sub-problem or arising from their interaction: they divide into hard and soft constraints, the former being mandatory and the latter being subject to penalties when violated. For practical purposes, we will refer to problem solutions where all hard constraints are satisfied as a *stable*, *unstable* otherwise.

Given a test instance, all the information required are provided, including pre-admitted occupants, patients, nurses, surgeons, operating theaters, rooms, shifts, scheduling dimension in days and penalty weights for soft constraints.

2 Approach

The proposed approach involves the application of the *Tabu Search* (TS) algorithm, which is it is a local search method that aims to find the global optimum of a combinatorial optimization problem. [2]

The main idea behind TS is to enhance local search performances by relaxing its rules: in particular, worsening moves can be accepted if no improving move is available, which prevent from getting stuck in local optima, and previously-visited solutions are discouraged, which prevent cycles. This last aspect is implemented storing moves in a queue, called *tabu list*, the dimension of which needs to be tuned. In order to mitigate the constraints introduced by the list, the algorithm also includes the possibility to accept forbidden moves when they met a certain *aspiration criterion*, for example if they lead to a better solution than the current one.

The next paragraphs will discuss the main aspects that need to be addressed in order to adapt TS to IHTP.

2.1 Solution Representation

The current status of the problem is represented by three main data structures:

- *Patient Admission and Scheduling (PAS)*: boolean array of size $n_{\text{days}} \times n_{\text{rooms}} \times n_{\text{patients}}$. For each patient, it indicates which room they are assigned to during all the days of recovery.
- *Surgical Case Planning (SCP)*: boolean array of size $n_{\text{days}} \times n_{\text{patients}} \times n_{\text{surgeons}} \times n_{\text{operating theaters}}$. For each patient, it indicates the day of admission (i.e. surgery) along with the predetermined surgeon and the chosen operating theater.

- *Nurse-to-Room Assignment (NRA)*: boolean array of size $n_{\text{shifts}} \times n_{\text{rooms}} \times n_{\text{nurses}}$. For each nurse, it indicates which room they are assigned to during each shift.

In order to lower computational time of certain soft constraints, two other structures are used: *Workload* and *Skill level*, both arrays of size $n_{\text{shifts}} \times n_{\text{rooms}} \times n_{\text{patients}}$, storing the workload and skill level required by a patient assigned to a room during a certain shift.

2.2 Starting Point

As already mentioned, a test instance starts with some patients already admitted to the hospital, called occupants. This configuration clearly is an unstable state of the problem, since no hard constraint is guaranteed to be satisfied. With these premises, two approaches can be considered in order to generate a starting point: create a stable configuration from scratch using advanced and exhaustive combinatorial techniques, which is a very complex task both in terms of time and space complexity and would partly defy the purpose of this work; otherwise, accept the unstable configuration as it is and let TS find a stable solution by opportunely tuning its behavior. The second approach is the one that has been chosen, since it is more efficient and allows to fully exploit the potential of the algorithm. Basically, the main idea is to generate a neighborhood that is able to direct the algorithm towards a stable status, and hopefully the algorithm will be able to find a final solution that satisfies all the hard constraints. The choices taken to accomplish this non-trivial task are presented in the next paragraphs.

2.3 Neighborhood Definition

Given the nature of the problem and the constraints involved, the type of moves performed by the algorithm are the following:

- Patient
 - Scheduling: for a patient, admit them to a room on a day and establish the operating theater.
 - Unscheduling: for a patient, remove them completely from the schedule.
- Nurse
 - Assignment: for a nurse, assign them to a room during a shift.
 - Unassignment: for a nurse, remove them from an assigned room during a shift.

The neighborhood is defined as the set of all possible moves that can be performed on the current solution, meaning that for each of the aforementioned types of moves all the possible combinations of patients, rooms, days, operating theaters, nurses and shifts are considered. Obviously, moves effectively considered are those that do not violate hard constraints and do not destabilize the solution. Moreover, given that the current status of the problem is not guaranteed to be stable, the neighborhood needs to be ulteriorly filtered if necessary: more precisely, when there are still non-admitted mandatory patients, these have priority over the others and scheduling regarding optional patients is ignored.

Since TS operates comparing the current moves with the ones stored in the tabu list, it is also necessary to define a way to compare two moves. For patient moves, the comparison is performed by checking the type and all the parameters involved (i.e. patient, room, day, operating theater), meaning that only the same exact action is considered tabu. For nurse moves, on the other hand, equality is asymmetric and comparison follows these rules: when considering an assignment, it can only be equal to an assignment with the same parameters (i.e. nurse, room, shift), so an exact copy; when considering an unassignment, it can be equal to both an assignment or unassignment with equal nurse and room, independently of the shift. The reason behind this choice is that the nurses assignment is a vital part of the schedule, since room coverage needs to be granted in order to admit patients: the strategy is to force the algorithm in a direction where nurses are more likely to be assigned and avoids the situation where nurses are assigned and unassigned multiple times consecutively just because the shift changes.

As a consequence of all these choices, we expect the optimization to be more stable for hard constraint satisfaction, specifically mandatory patients admission and room coverage, but less effective in minimizing soft constraints focusing on nurses.

2.4 Tabu List Size

The tabu list size is a crucial parameter of the algorithm, depending both on the problem dimension (i.e. number of patients, nurses, rooms, etc.) and the neighborhood definition we established. A small tabu list size can lead

to cycling and getting stuck to a local space, while a large one can slow down the algorithm and make it less effective.

In our case, TS evaluates moves and performances based on a the penalty of the soft constraints, moreover it has been adapted to consider the hard constraints as well. As a consequence, the two previous edge cases translates as follows:

- Small tabu list: the algorithm tends to reaccept previous moves and performs a small set of actions, more precisely nurses assignment and unassignment. In other words, TS focus more on penalty optimization and get stuck in a local space where the best action at each step is to assign or unassign a nurse, which slightly improves the penalty or keeps it constant. This is not desirable, since the hard constraints are not guaranteed to be satisfied.
- Large tabu list: the algorithm tends to avoid previous moves and explores a larger space, performing patient scheduling. However, patient unschedule actions remain tabu for a long time preventing TS to free up a room, which means the current nurse assigned to it is stuck and cannot be exchanged with a better nurse in terms of soft constraints. This is not desirable as well, since the penalty is not minimized.

These cases are the extremes of the spectrum, but evidence well how the tabu list size is the key parameter in order to opportunely balance the algorithm and avoid having hard constraints violated, nor them prevailing over the soft ones.

2.5 Aspiration Criterion

In this case, the aspiration criterion is used to accept forbidden moves when they lead to a better solution than the best overall one found so far.

3 Results

The algorithm was tested with some provided instances, whose parameters are shown in Table 1 along with the chose size for the tabu list. In all the instances, the algorithm was able to reach a solution where no hard constraint was violated, with penalty values shown in Table 2.

Instance	Days	Shifts	Occupants	Patients	Operating Theaters	Rooms	Nurses	Tabu List Size
toy	7	3	2	7	1	3	11	30
1	21	3	7	42	2	5	13	150
2	14	3	5	37	4	6	17	200
3	14	3	10	45	2	6	14	200
4	14	3	7	54	3	8	19	250
5	14	3	9	62	2	6	15	200
6	14	3	8	110	3	9	20	300

Table 1: Test instances parameters and tabu list size.

Test Instance	Starting penalty	Tabu Penalty	Proposed Penalty
toy	2100	269	292
1	6350	3976	3177
2	12950	2443	1583
3	15758	11624	10184
4	13501	6524	2332
5	24805	16275	15713
6	55500	25366	17048

Table 2: Penalty reached by the Tabu Search algorithm on the test instances, compared to the penalties of the starting configuration and the competition proposed solution.


4 Discussion

The proposed approach performances are overall satisfactory, with the algorithm being able to find a solution guaranteeing no hard constraint violation in all the tested instances. The penalty values reached by TS are sometimes competitive with the precomputed solutions, but are generally higher, especially for the larger problems.


Focusing on the first instances, we notice that they achieve great results in term of patient scheduling, with an high percentage of penalty due to non minimization of nurses soft constraints like continuity of care. As mentioned before, this happens when the tabu list size is too large: however, in this cases, tests performed with smaller tabu list size did not lead to better results, but to a higher penalty value or even violation of hard constraints. One possible solution to this issue could be to implement the possibility for nurse substitution, without the need to free up the room and remove the assigned nurse beforehand: this would allow the algorithm to maintain stability and optimize nurse assignment in a more effective way. It is worth mentioning, however, that this would require a more complex neighborhood definition and a more sophisticated tabu list management, since the algorithm would need to consider the possibility of swapping nurses between rooms and shifts. As a matter of fact, a bad management of this feature could lead to the algorithm cycling between nurses interchange, significantly worsening the performances. For this reasons, an implementation of this kind would surely translate in a more complex and slower algorithm.

A Appendix


A.1 Plots




(a) Toy.




(b) Test 01.




(c) Test 02.




(d) Test 03.



(e) Test 04.



(f) Test 05.



(g) Test 06.

Figure 1: Costs of the solution found by the Tabu Search algorithm on the test instances.

References

- [1] URL: <https://ihtc2024.github.io/>.
- [2] Fred Glover. “Tabu Search—Part I”. In: *ORSA Journal on Computing* 1.190-206 (3 1989). URL: <https://doi.org/10.1287/ijoc.1.3.190>.